

Implementing Improvements for Secure Function Evaluation

A Thesis
Presented to
The Division of Mathematics and Natural Sciences
Reed College

In Partial Fulfillment
of the Requirements for the Degree
Bachelor of Arts

Alex Ledger

May 2016

Approved for the Division
(Mathematics)

Adam Groce

Acknowledgements

I want to thank a few people.

Preface

This is an example of a thesis setup to use the reed thesis document class.

List of Abbreviations

You can always change the way your abbreviations are formatted. Play around with it yourself, use tables, or come to CUS if you'd like to change the way it looks. You can also completely remove this chapter if you have no need for a list of abbreviations. Here is an example of what this could look like:

{AL-0: Do I need this?}	ABC	American Broadcasting Company
	CBS	Columbia Broadcasting System
	CDC	Center for Disease Control
	CIA	Central Intelligence Agency
	CLBR	Center for Life Beyond Reed
	CUS	Computer User Services
	FBI	Federal Bureau of Investigation
	NBC	National Broadcasting Corporation

Table of Contents

Introduction	1
Chapter 1: Cryptographic Primitives	3
1.1 Introducing Cryptographic Security	3
1.2 Encryption	6
1.3 Computational Indistinguishability	9
1.4 Boolean Circuit	10
1.5 Oblivious Transfer	11
Chapter 2: Classic 2PC	13
2.1 2PC Security Motivation	13
2.2 2PC Security Definition	16
2.3 Yao's Garbled Circuit	19
2.3.1 Security of Garbled Circuits	22
2.3.2 Notes about complexity	25
2.4 GMW	26
Chapter 3: Improving MPC	27
3.1 Point and Permute	29
3.2 Garbled Row Reduction 3	31
3.3 Free XOR	32
3.4 Garbled Row Reduction 2	34
3.5 FleXOR	35
3.6 Half Gates	37
3.7 Improving Oblivious Transfer	39
Chapter 4: Chaining Garbled Circuits	41
4.1 Chaining	42
4.2 Security of Chaining	44
4.3 Single Communication Multiple Connections	44
4.4 Security of SCMC	45
Chapter 5: Implementation	47
5.1 CompGC	47
5.2 LibGarble and JustGarble	49

5.3	Our Implementation: CompGC	49
5.4	Adding SCMC	49
5.5	Experimental Setup	49
5.6	Experiments	50
5.7	Results	50
Chapter 6: Future and Related Work		51
Conclusion		53
Appendix A: The First Appendix		55
Appendix B: The Second Appendix, for Fun		57
References		59

List of Tables

1.1	The mapping of an XOR gate.	10
1.2	The truth table of the less than circuit.	12
2.1	A garbled table for an AND gate with input wires W_0 and W_1 and output wire W_5	21
3.1	Summary of Garbled Circuit Improvements. GRR3 stands for garbled row reduction 3 and GRR2 stands for garbled row reduction 2. {AL-24: two halves make a whole paper}	29
3.2	Garbled Gate for Point and Permute	31
3.3	Example garbled gate using point and permute and garbled row reduction 3. The gate being computed is given in figure Make it 23:30 in mike's talk	32
3.4	Left: garbler-half-gate garbled table for $a = 0$. Middle: garbler-half-gate garbled table for $a = 1$. Right: garbler-half-gate garbled table for arbitrary a	38
3.5	Evaluator-half-gate garbled table.	38

List of Figures

1.1	An AND gate.	11
1.2	An XOR gate.	11
1.3	A NOT gate.	11
1.4	A circuit that computes the less or equal to function, equivalent to the less than function for input of two one-bit values.{AL-8: make with tikz}	12
2.1	A high level overview of the garbled circuit protocol	19
2.2	A simple boolean circuit. {AL-12: Draw gate and wire number notation in here}	19

Abstract

The preface pretty much says it all.

Dedication

You can have a dedication here if you wish.

Introduction

{AL-4: in primitives, should I clarify what I mean by "break an algorithm"?} {AL-5: Needs to be redone to emphasize chaining, and the work that I did. Right now, it just generally motivates SFE}

Secure Function Evaluation (SFE) is the study and creation of protocols for securely computing a function between multiple parties. Secure, in this context, means that the protocol prevents each party from learning anything about the other parties.

The idea is best communicated through an example: suppose Alice and Bob are millionaires and wish to determine who is wealthier. But Alice and Bob are also secretive, and do not want to disclose their exact amount of wealth. Is there some method by which they can determine who has more money?

{AL-6: check if using greater than or less than function} Alice and Bob can solve their problem by specifying a function f ; in this case f takes two inputs x and y and outputs 1 if $x < y$, and 0 otherwise (f is the less than function). If Alice and Bob can secretly give f their amount of wealth and then retrieve the answer from f , then their problem is solved: they have securely computed who is wealthier.

The goal of SFE broadly is to give Alice, Bob and their friends the ability to securely compute an arbitrary function at their will. For the protocol to be secure, it needs to have the following informal properties:

- **Privacy:** Each party's input is kept secret.
- **Correctness:** The correct answer to the computation is computed.

Originally, the goal of the research community was to develop a secure protocol, define security, and prove the protocol is secure. In more recent times, the focus has shifted to making SFE protocols faster.

If SFE could be made fast enough, it could serve a wide range of applications. For example, imagine that two companies who operate in a similar industry want to work together, but they don't want to disclose any company research which the other doesn't know. These companies could run a set intersection function (a function that given two inputs finds their intersection, or overlap), to determine what information they can disclose without giving away important information.

Another example is running algorithms on confidential medical data. Consider the case where various hospitals want to jointly run algorithms on their data for medical research. The hospital needs guarantees, for both legal and ethical reasons, that the data will not be disclosed as the algorithm is being run, as medical data can be very sensitive. The hospitals could use SFE to run computation on all of their data, maintaining the privacy of their patients.

Since research into SFE has focused on creating a method by which an arbitrary function can be computed securely, the application of SFE may be beyond what we can presently conceive of. It's not unlikely that SFE will become a standard in the internet, where when you access the internet, behind the scenes your access is being plugged into an SFE protocol, sent off to another computer to do some processing. For this future to be realized, work needs to be done to make SFE faster and more flexible.

Chapter 1

Cryptographic Primitives

Secure Computation is the study of computing functions in a secure fashion. SC is split into two cases: cases where there are two parties involved, referred to as two-party computation (2PC), and cases where there are three or more parties, referred to as multiparty computation (MPC). This thesis will focus primarily on 2PC protocols, but many of the methods are applicable to MPC as well.

2PC protocols are complex cryptographic protocols that rely on a number of cryptographic primitives. In order to understand 2PC, it is not crucial to understand how the cryptographic primitives work, but it is important to understand their inputs, outputs and security guarantees. This first chapter will give an overview of the cryptographic primitives used in 2PC protocols.

1.1 Introducing Cryptographic Security

The goal of this section is to explain cryptographic security, starting at an intuitive level and moving outward. We do not present cryptographic security in a comprehensive fashion; rather, we explain cryptographic security with the goal of explaining 2PC protocols and their security. For more information on cryptographic security, we encourage the reader to peruse [Katz & Lindell \(2007\)](#).

We define a few intuitive terms to get started. A *cryptographic scheme* is a series of instructions designed to perform a specific task. An *adversary* is an algorithm that tries to *break* the scheme. If an adversary *breaks* a scheme, then the adversary has learned information about inputs to the scheme that they shouldn't.

The goal of defining security in cryptography is to build a formal definition that matches real world needs and intuitions. A good starting place is to consider perfect security. A scheme is perfectly secure if no matter what the adversary does, they cannot break the scheme. Even if the adversary has unlimited computational power, in terms of time and space, an adversary cannot break a perfectly secure scheme.

However, perfect security is not the most useful way to think of security, because it makes forces schemes to be slow and communication intensive. We relax the definition of security by requiring that the adversary run in polynomial time. This substantially reduces the power of the adversary, and it matches our intuition. We are really only concerned with what adversaries can reasonably achieve, as opposed to theoretically possible.

Because computers have improved drastically over the years, what was previously considered a reasonable adversary is not what is considered a reasonable adversary today. Modern computing advances have created easy access to faster computation, meaning that modern adversaries can solve harder problems than they could in previous years. As a concrete example, consider an giving an adversary the following problem: find the factors of N . The average computer today can solve the problem for a larger N than the average computer a decade ago.

Changing computational power makes it important to scale how hard a cryptographic scheme is to break. To this end, we introduce a *security parameter*, denoted as λ . A security parameter is a positive integer that represents how hard a scheme is to break. A larger security parameter should mean that the scheme is more difficult to break. More specifically, the security parameter is correlates to the input-size

of the problem underlying the cryptographic scheme. For example, if the underlying problem is factoring large number N , then $N = \lambda$. As N and λ scale up, the factoring problem becomes more difficult and the scheme becomes hard to break.

Finally, we acknowledge that adversaries have access to some random values, hence we strengthen adversaries to be probabilistic algorithms. Probabilistic means that the algorithms have access to a string of uniform random bits, with the implication being that the algorithm is capable of guessing.

In order to reason about the security of cryptographic schemes, it is useful to think about breaking a scheme in terms of a probability. For example, we want to be able to say that the best adversary, that is best probabilistic polynomial-time algorithm, has some probability p of breaking the scheme. We note that p is nonzero, since the adversary can always guess and be right with some nonzero probability. To achieve a probability based formalism, we introduce a negligible function. Informally, a negligible function is smaller than the reciprocal of all polynomial functions. Formally, a negligible function is:

Definition 1 A function $\mu : \mathbb{N} \rightarrow \mathbb{R}$ is negligible if for all positive polynomial $p(\cdot)$, there exists positive integer N_p such that for all $x > N_p$,

$$|\mu(x)| < \frac{1}{p(x)}. \quad (1.1)$$

Goldreich (1995)

◇

Examples of negligible functions include 2^{-n} , $2^{-\sqrt{n}}$ and $n^{-\log n}$.

To put a negligible function to use, say an adversary is attacking a cryptographic scheme that is equivalent to solving a problem P with input-size λ and 2^λ possible answers. Moreover, say that P is known to be NP-hard such that there is no polynomial time algorithm to solve P . Then, the best that the adversary can do is to guess the answer to P . Hence the probability that the adversary that finds the answer the

P , or breaks the scheme, is

$$\Pr[A \text{ correctly answers } P] = 2^{-\lambda}$$

Since $2^{-\lambda}$ is a negligible function, we say that the adversary has a negligible probability of breaking the scheme.

In summary, we model an adversary as a probabilistic polynomial-time algorithm. This limits the computational power of the adversary to what is reasonably computable in reality. Moreover, we can scale the security of a scheme or problem by changing λ , the security parameter. A higher security parameter makes the scheme more difficult to break.

1.2 Encryption

Encryption is the process of obfuscating a message, and then later un-obfuscating the message. Say Alice has a message that she wants to send to Bob, but somewhere between Alice and Bob sits Eve, who wants to learn about the message. An encryption scheme enables Alice to send her message to Bob with confidence that Eve cannot learn any information about the message.

An encryption scheme is composed of three algorithms: **Enc**, **Enc**⁻¹ and **Gen**; formally, we say an encryption scheme is a tuple $\Pi = (\text{Gen}, \text{Enc}, \text{Enc}^{-1})$ ¹. **Enc** the obfuscating algorithm, **Enc**⁻¹ is the un-obfuscating algorithm and **Gen** generates a key. The key is extra information that **Enc** and **Enc**⁻¹ use to obfuscate and un-obfuscate the message respectively. The key, denoted k , is a random² string of λ bits, that is k is randomly sampled from $\{0, 1\}^\lambda$ where λ is the security parameter of the

¹We use Π here to denote the encryption scheme, because it is a protocol. Protocol starts with a p.

²The notion of randomness in cryptography is precisely defined, and in cases where λ is large, it is sufficient for k to be pseudorandom. Pseudorandomness is also precisely defined.

encryption scheme. As per the discussion on security parameters, as λ increases and k grows in length, an encryption scheme should become harder to break.

Enc, the encryption algorithm, takes a message and the key as input and outputs an obfuscated message. Enc^{-1} , the decryption algorithm, takes the encrypted message and the key as input and outputs the original message. We refer to the original message as the plaintext or pt and the encrypted message as the ciphertext or ct .

$$\begin{aligned}\text{Gen}(1^n) &\rightarrow k \\ \text{Enc}_k(pt) &\rightarrow ct \\ \text{Enc}_k^{-1}(ct) &\rightarrow pt\end{aligned}\tag{1.2}$$

We are not concerned with how encryption schemes are implemented or on what problems they rely; rather, we use encryption schemes as subroutines, so we are concerned with the security guarantees that they offer.

We say that an encryption scheme is secure if an adversary cannot tell the difference between two messages. We define security using a thought experiment. In the thought experiment, the adversary has access to the encryption algorithm with key hardcoded in. This means that the adversary can encrypt any message they want, and see how the message would encrypt. The goal of the adversary at this point in the thought experiment is to find a pattern or weakness in the encryption algorithm that they can exploit. The adversary eventually picks any two messages m_0 and m_1 which they did not give to their encryption algorithm and see how they encrypt. The adversary shows us m_0 and m_1 . We choose one of the messages³, encrypt the message, and send the resulting ciphertext to the adversary. The adversary's goal now is to determine which message we encrypted. They still may use their encryption algorithm with the key hardcoded in. Eventually the adversary must output either 0 or

³We select the message uniformly at random.

1 indicating that they think we selected m_0 or m_1 respectively. If the adversary picks correctly, then we say that the adversary wins; otherwise, we say that the adversary loses.

Finally and informally, the encryption scheme is considered secure if the probability that the adversary wins is $\frac{1}{2} + \mu(\lambda)$, i.e. the best the adversary can do is guess.

Definition 2 An encryption scheme is secure under a chosen-plaintext attack if for all probabilistic polynomial-time adversaries \mathcal{A} , there exist a negligible function μ such that

$$\Pr[E_{\mathcal{A},\Pi}(n) = 1] \leq \frac{1}{2} + \mu(n) \quad (1.3)$$

where E is the following experiment:

1. Generate key k by running $\text{Gen}(1^n)$.
2. The adversary \mathcal{A} is given 1^n and oracle access to Enc_k , and outputs a pair of messages m_0 and m_1 of the same length.
3. A uniform bit $b \in \{0, 1\}$ is selected, and then a ciphertext $c \leftarrow \text{Enc}_k(m_b)$ is computed and given to \mathcal{A} .
4. \mathcal{A} continues to have oracle access to Enc_k , and outputs a bit b' .
5. The output of the experiment is defined to be 1 if $b' = b$ and 0 otherwise. In the former case, we say that \mathcal{A} succeeds.

◇

It is useful for 2PC to create an encryption scheme that requires two keys to encryption and decrypt. An encryption scheme with two keys is called a *dual-key cipher* (DKC) [Bellare et al. \(2012\)](#). It is easy to instantiate a DKC if one has a secure single-key encryption scheme: let k_0 and k_1 be two keys and instantiate the

DKC as follows:

$$\begin{aligned} \text{EncDKC}_{k_0, k_1}(pt) &= \text{Enc}_{k_1}(\text{Enc}_{k_0}(pt)) \\ \text{EncDKC}_{k_0, k_1}^{-1}(ct) &= \text{Enc}_{k_0}^{-1}(\text{Enc}_{k_1}^{-1}(ct)) \end{aligned} \tag{1.4}$$

If the encryption scheme used to create the DKC is secure, then it is easy to see that the DKC is also secure. We are formally considering the statement: $\text{Enc secure} \implies \text{DKC secure}$. Consider the contrapositive: $\text{DKC insecure} \implies \text{Enc insecure}$. If the DKC is insecure, then an adversary can find two messages, m_0 and m_1 such that their ciphertexts are distinguishable. {AL-7: Figure this out, need to break the old intro crypto work}

1.3 Computational Indistinguishability

This section introduces the idea of computational indistinguishability. We do not use computational indistinguishability immediately, but it will be important later for defining security of a 2PC protocol.

Informally, two probability distributions are indistinguishable if no polynomial-time algorithm can tell them apart. The thought experiment is like this: an algorithm is given one of two distributions. If the algorithm correctly determines which distribution it was given, then it wins, otherwise the algorithm loses. The algorithm, since it must run in polynomial time, can only sample a polynomial number of values from the distributions.

Formally, computational indistinguishability is:

Definition 3 Let $\mathcal{X} = \{X_n\}_{n \in \mathbb{N}}$ and $\mathcal{Y} = \{Y_n\}_{n \in \mathbb{N}}$ be distribution ensembles. \mathcal{X} and \mathcal{Y} are computationally indistinguishable, denoted $\mathcal{X} \approx_C \mathcal{Y}$, if for all probabilistic

polynomial-time algorithms D , there exists a negligible function μ such that:

$$|Pr_{x \leftarrow X_n}[D(1^n, x) = 1] - Pr_{y \leftarrow Y_n}[D(1^n, y) = 1]| < \mu(n) \quad (1.5)$$

Katz & Lindell (2007).

◇

We quickly break down the definition. The unary input 1^n tells the algorithm D to run in polynomial time in n . The probability distributions X_n and Y_n are restricted by n , which in this context is the security parameter. The phrases $x \leftarrow X_n$ and $y \leftarrow Y_n$ mean that the probability is taken over samples from the distributions.

1.4 Boolean Circuit

A function in a 2PC protocol is represented as a boolean circuit. A boolean circuit takes as input $x \in \{0, 1\}^n$, performs a series of small operations on the inputs, and outputs $y \in \{0, 1\}^m$. You may have encountered circuits and logical operators in another context, where the inputs and outputs were True and False. For our usage, True will correspond to the value 1, and False will correspond to the value 0.

The small operations done inside of a circuit are performed by a *gate*. A gate is composed of three wires: two input wires and one output wire, where a *wire* can have a value either 0 or 1. A gate performs a simpler operation on the two inputs, resulting in a single output bit. Table 1.4 gives the mapping of an XOR gate.

x	y	xor(x,y)
1	1	0
1	0	1
0	1	1
0	0	0

Table 1.1: The mapping of an XOR gate.

A circuit is a combination of gates that are stringed together. It turns out that circuits are quite powerful: in fact, a circuit composed only of AND gates, XOR

gates and NOT gates can compute any function or algorithm f . In other words, if there's some algorithm that can do it, then there is some circuit that can do it as well. Figure 1.4 shows the circuit representation of the less than function, f as specified in equation ??.



Figure 1.1: An AND gate.



Figure 1.2: An XOR gate.



Figure 1.3: A NOT gate.

1.5 Oblivious Transfer

Oblivious Transfer (OT) is a communicating protocol that underlies many more complicated cryptosystems. At the highest level, Alice potentially sends two messages to Bob, and Bob only receives one of the messages.

Consider the following thought experiment: Alice and Bob want to exchange a message. Alice has two messages m_0 and m_1 , and she Alice wants to send one of them Bob but does not care which one. Bob knows that he wants message m_b where $b \in \{0, 1\}$. Alice gives both messages to the trusted post-office. Bob also goes to the trusted post office, and says that he wants m_b . The post office gives m_b to Bob, throws the other message (m_{1-b}) in the trash, and does not tell Alice which message



Figure 1.4: A circuit that computes the less or equal to function, equivalent to the less than function for input of two one-bit values. [{AL-9: make with tikz}](#)

x	y	$x < y$
0	0	0
0	1	1
1	0	0
1	1	0

Table 1.2: The truth table of the less than circuit.

they gave to Bob. Now, Alice has sent a message to Bob, but is oblivious as to which message was sent.

Oblivious Transfer is a useful protocol and is a fundamental component of 2PC protocols. Oblivious Transfer protocols guarantee the following security properties:

1. Alice does not know whether Bob recieved m_0 or m_1 .
2. Bob does not know anything about m_{1-b} , the message that he did not receive.

Chapter 2

Classic 2PC

[{AL-10: this paragraph is a little weak}](#) Secure Computation was first proposed in an oral presentation by Andrew Yao [?]. Since Yao's presentation, two main methods have emerged for secure computation. The first method is garbled circuits, and is the focus of this chapter. The second method is GMW, and is explained in brief at the end of the chapter.

This chapter begins with motivating and describing desirable properties of a 2PC protocols, culminating in a definition of security. Next, the chapter describes garbled circuits, a popular method for performing 2PC.

2.1 2PC Security Motivation

Think back to Alice and Bob from the introduction. Alice and Bob are millionaires who wish to determine who is wealthier without disclosing how much wealth they have. More formally, Alice has input x and Bob has input y (x and y are integers corresponding to the wealth of each party), and they wish to compute the less than function f , such that

$$f(x, y) = \begin{cases} 1 & \text{if } x < y; \\ 0 & \text{otherwise.} \end{cases} \quad (2.1)$$

We call the overarching interaction between Alice and Bob protocol Π , and Π consists of all messages exchanged and computations performed. Based on the setup of the problem, we can list a few properties that Alice and Bob wish Π to have.

Privacy Parties only learn their output. Any information learned by a single party during the execution of Π must be deduced from the output. For example, if Alice learns that she had more money after computing f , then she learns that $y < x$; however, this information about y is deducible from the output therefore it is reasonable. It would be unreasonable if Alice learns that $1,000,000 < y < 2,000,000$, as that information is not deducible from $f(x, y)$.

Correctness Each party receives the correct output. In the case of Alice and Bob, this simply means that they learn correctly who has more money. In particular, correctness means that Alice and Bob *both* learn the output.

One possible method for constructing a definition of security would be to list a number of properties a secure protocol must have. This approach is unsatisfactory for a number of reasons.

One reason is that an important security property that is only relevant in certain cases may be missed. There are many applications of 2PC, and in some cases, there may be certain properties that critical to security. Ideally, a good definition of 2PC works for all applications, hence capturing all desirable properties. A second reason that the property based definition is unsatisfactory is that the definition should be simple. If the definition is simple, then it should be clear that *all* possible attacks against the protocol are prevented by the definition. A definition based on properties in this respect as it becomes the burden of the prover of security to show that all relevant properties are covered [Lindell & Pinkas \(2009\)](#).

We must also think about the aims of each party involved in the protocol. Can we trust that parties are going to obey the protocol? Are the parties going to try to

cheat? These considerations are called the *security setting*. There are two primary security settings: the semi-honest setting and the malicious setting.

The work presented in this thesis uses the semi-honest setting. In the semi-honest setting, we assume that each party obeys the protocol but tries to learn as much as possible from the information they are given. This means that parties do not lie about their information, they do not abort, they do not send or withhold messages out of order, or deviate in any way from what is specified in the protocol. In contrast, the malicious setting considers that each party is liable to lie and cheat; parties can take any action to learn more information.

The malicious setting is much more realistic. Parties that are involved in cryptographic protocols are liable to lie and cheat, for why else would they even be engaged in the cryptographic protocol in the first place? There are two main reasons why the semi-honest setting is useful. The first is that many protocols can be constructed for the semi-honest setting, and then improved to function in the malicious setting. There is a strong history of this occurring with protocols. It's simply easier to think through and create protocols for the semi-honest setting; at the very least, it's a valuable starting point for building complex cryptosystems. In the case of 2PC, there exist malicious protocols, and in fact, the primary 2PC protocol that this thesis uses, garbled circuits, can be improved to be malicious secure without too much difficulty.

The second reason that the semi-honest setting is that it does have use cases in the real world. There are some scenarios where parties want to compute a function amongst themselves, and trust each other to act semi-honestly. One example is hospitals sharing medical data. Hospitals are legally, and arguably ethically, restricted from sharing medical data, but this data can have great value especially when aggregated with datasets from other hospitals. 2PC offers hospitals the means to “share” their data, perform statistics and other operations on it, while keeping the data entirely private. Other examples where semi-honesty is sufficient include mutually trusting

companies and government agencies.

2.2 2PC Security Definition

In this section we discuss at a high level the definition of 2PC security, and then give the formal definition. The definition of security for 2PC protocols is the most complicated cryptographic theory that we have encountered thus far.

Recall our setup: Alice and Bob are semi-honest parties with inputs x and y respectively who wish to compute the function $f : \{0, 1\}^n \times \{0, 1\}^n \rightarrow \{0, 1\}^m \times \{0, 1\}^m$. The protocol Π is a 2-party computation protocol, that is instructions that Alice and Bob follow, that enables them to compute f . To think about the security of Π , we imagine an ideal world where the protocol is computed securely, and compare the ideal world to the real world where Π is executed.

In the ideal world, we imagine that there is a third, trusted and honest party Carlo. Instead of Alice and Bob communicating amongst each other, Alice and Bob send their inputs to Carlo. Carlo computes $f(x, y)$ himself, and sends the output to Alice and Bob. The only information that Alice and Bob have in the ideal world is their individual inputs and the output.

Informally, we say that Π is secure if Alice and Bob learn *essentially the same* information executing Π in the real world as they do computing f in the ideal world with Carlo. If either Alice or Bob manage to learn more information in the real world, then the protocol Π is leaking information that cannot be deduced from their individual inputs and the output of f . This idea of security is formally achieved using the concept of computational indistinguishability presented in Chapter 1. Recall that computational indistinguishability is the idea that two probability distributions are essentially the same - no polynomial time algorithm can distinguish them. To use computational indistinguishability, we think of the information that Alice and Bob

learn in the real and ideal world as probability distributions. The idea may seem fuzzy now, but the idea will become clearer as the probability distributions are explained

To construct the probability distribution of the ideal world, we introduce *simulators* S_A and S_B . S_A and S_B are probabilistic polynomial-time algorithms who are essentially adversaries, like the adversary in the definition of encryption from Chapter 1, that specifically attack the ideal world. Simulator S_A takes as input x , Alice's input, and $f(x, y)$, the output of the function because that is the information that Alice has access to in the ideal world. Likewise, S_B takes input y and $f(x, y)$, since that is the information that Bob has access to in the ideal world.

To create a probability distribution, we consider what S_A does over all possible input: the distribution of the possible outputs is given by $\{S_A(x, f(x, y))\}_{x \in \{0,1\}^*}$. Let us break this distribution down: S_A is a fixed single algorithm. x is Alice's input, and $f(x, y)$ is the output of the function. The set is indexed by all possible x , so all possible inputs that Alice could have. In summary, $\{S_A(x, f(x, y))\}_{x \in \{0,1\}^*}$ represents the possible information that an algorithm could deduce from all possible x and $f(x, y)$. We think of $\{S_B(y, f(x, y))\}_{y \in \{0,1\}^*}$ for Bob's input similarly.

In the real world, we need to consider what information Alice and Bob have at their disposal. Recall that Alice and Bob are semi-honest, which means that Alice and Bob follow all instructions of Π , but they use any information they receive along the way. More precisely, Alice and Bob obey the protocol, but also maintain a record of all intermediate computations. We call Alice's record of intermediate computations Alice's *view*, $\mathbf{view}_A(x, y)$, which depends on inputs x and y . And now we create the probability distribution for Alice: $\{\mathbf{view}_A(x, y)\}_{x, y \in \{0,1\}^*}$. This distribution represents the Alice's information from the intermediate computation indexed over all possible inputs x and y . Likewise, we call Bob's record of intermediate computations Bob's *view*, $\mathbf{view}_B(x, y)$, and his probability distribution of intermediate computations is $\{\mathbf{view}_B(x, y)\}_{x, y \in \{0,1\}^*}$.

To wrap it up, if Π in the real world is the same as the Alice, Bob and Carlo in the ideal world, then the simulator S_A and S_B should only be able to learn what can be learned from the intermediate computations. That is, the probability distributions $\{S_A(x, f(x, y))\}_{x,y \in \{0,1\}^*}$ and $\{\text{view}_A(x, y)\}_{x,y \in \{0,1\}^*}$ should be essentially the same, i.e., computationally indistinguishable.

With this intuition in mind, we give Goldreich's definition of 2PC security from his textbook *Foundations of Cryptography Volume II* [Goldreich \(1995\)](#).

Definition 4 Let $f = (f_1, f_2)$ be a probabilistic, polynomial time functionality where Alice and Bob compute $f_1, f_2 : \{0, 1\}^n \rightarrow \{0, 1\}^m$ respectively. Let Π be a two party protocol for computing f . Define $\text{view}_i^\Pi(n, x, y)$ (for $i \in \{1, 2\}$) as the view of the i th party on input (x, y) and security parameter n . $\text{view}_i^\Pi(n, x, y)$ equals the tuple $(1^n, x, r^i, m_1^i, \dots, m_t^i)$, where r^i is the contents of the i th party's internal random tape, and m_j^i is the j th message that the i th party received. Define $\text{output}_i^\Pi(n, x, y)$ as the output of the i th party on input (x, y) and security parameter n . Also denote $\text{output}^\Pi(n, x, y) = (\text{output}_1^\Pi(n, x, y), \text{output}_2^\Pi(n, x, y))$. Note that view_i^Π and output_i^Π are random variables whose probabilities are taken over the random tapes of the two parties. Also note that for two party computation.

We say that Π securely computes f in the presence of static¹ semi-honest adversaries if there exist probabilistic polynomial time algorithms S_1 and S_2 such that for all $x, y \in \{0, 1\}^*$, where $|x| = |y|$, the following hold:

$$\{(S_1(x, f_1(x, y), f(x, y)))\}_{x,y} \approx_C \{(\text{view}_1^\Pi(x, y), \text{output}^\Pi(x, y))\}_{x,y} \quad (2.2)$$

$$\{(S_2(x, f_2(x, y), f(x, y)))\}_{x,y} \approx_C \{(\text{view}_2^\Pi(x, y), \text{output}^\Pi(x, y))\}_{x,y} \quad (2.3)$$

◇

The definition requires that $|x| = |y|$; however, this constraint can be overcome

¹{AL-11: TODO} Mention what static means

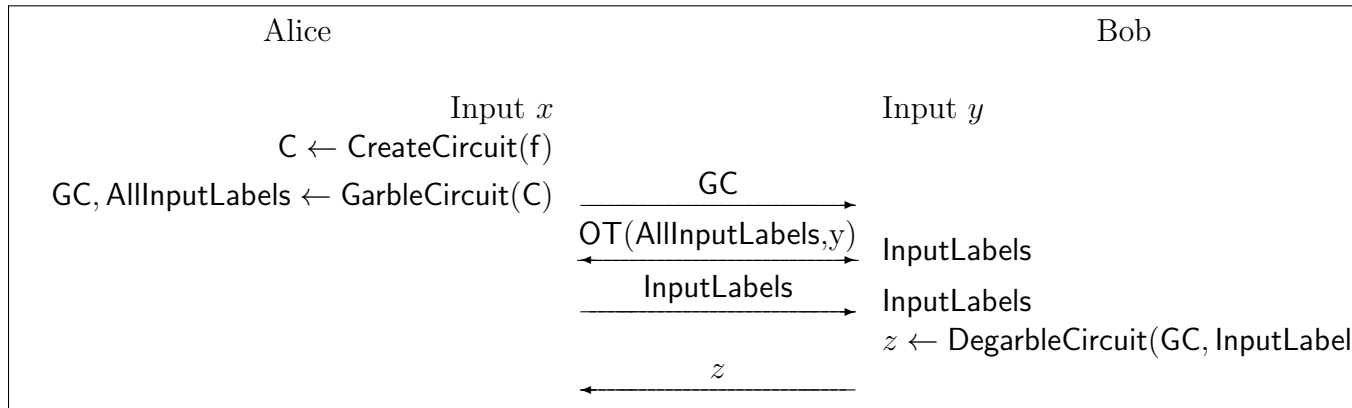


Figure 2.1: A high level overview of the garbled circuit protocol

by padding the shorter input.

A definition of 2PC security with malicious parties is substantially more complex. For more information on a malicious security definition, we refer the reader to ?.

2.3 Yao's Garbled Circuit

We now discuss a popular 2PC scheme called garbled circuits. At a high level, garbled circuits work by having one party, Alice, design a circuit that computes f . Alice encrypts (garble) that circuit and sends the encrypted circuit to Bob, along with some values corresponding to her and Bob's inputs. Bob decrypts the circuit, acquiring the output of $f(x, y)$ at the end.




Figure 2.2: A simple boolean circuit. {AL-13: Draw gate and wire number notation in here}

{AL-14: figure numbering is off} We now walk through how Alice garbles a circuit.

She starts with a typical boolean circuit, like the one shown in figure 2.3. In figure 2.3 the gates are ordered from 0 to 2, in order from nearest to the inputs to farthest from the inputs. Alice begins by assigning two wire labels to each wire - a wire is a line connecting the inputs and gates in the figure. Because Alice is working with a boolean circuit, each wire can be represented by either a 0 or 1. For a given wire W_i , the zeroth wire label W_i^0 represents 0, and the first wire label W_i^1 represents 1. We use W_i to represent the i th wire, W_i^j to represent the i wire's j label, and later we use W_i^* to represent one of wire i 's wire labels (without being specific about whether it is wire label 0 or wire label 1). A wire labels is a ciphertext, the output of the encryption algorithm². Alice assigns the wire labels by randomly sampling from $\{0, 1\}^\lambda$, where λ is the size of the output of the encryption algorithm.

After assigning wire labels, Alice first garbles gate G_0 , the gate nearest to the inputs³. Alice garbles by creating a garbled table T_{G_0} for G_0 . Gate G_0 is an AND gate, so the structure of the table resembles the logical table for an AND gate. The table has four columns. The first two columns are the input wire labels; in this case, these are wires W_0 and W_1 . The third column is the wire labels for the output wire; this case W_5 . The wire label placed in the third column is based on logical operation of the AND gate. For example, the third column and third row has the wire label associated with 0, since AND outputs zero on input of 0 and 1. The fourth column is the dual-key cipher encryption⁴ of the value in the third column, using the values of the first two columns as keys. The garbled table for G_0 is shown in table 2.3

Alice creates garbled tables all remaining gates; in this case, G_1 and G_2 remain. She then sends the fourth column all of garbled tables, the encryption of the respective out-wires, to Bob. Now if Bob has a label for each input wire, then Bob can acquire one of the labels of the output wire. More formally, say a gate was input wires W_i

²A common encryption algorithm used now is AES-128, so the wire labels are 128 bit strings.

³Multiple gates at some point could equidistant from the input. In these cases, the ordering of gates does not matter.

⁴See section {AL-15: what section?} for information on dual-key ciphers

W_0	W_1	W_5	Encryption
W_0^0	W_1^0	W_5^0	$\text{Enc}_{W_0^0, W_1^0}(W_5^0)$
W_0^1	W_1^0	W_5^0	$\text{Enc}_{W_0^1, W_1^0}(W_5^0)$
W_0^0	W_1^1	W_5^0	$\text{Enc}_{W_0^0, W_1^1}(W_5^0)$
W_0^1	W_1^1	W_5^1	$\text{Enc}_{W_0^1, W_1^1}(W_5^1)$

Table 2.1: A garbled table for an AND gate with input wires W_0 and W_1 and output wire W_5 .

and W_j and output wire W_k . Bob can acquire a wire label of W_k (i.e. W_k^0 or W_k^1) if he has one wire label of W_i (i.e. W_i^0 or W_i^1) and one wire label of W_j (i.e. W_j^0 or W_j^1).

In order to decrypt each gate, Bob needs to first acquire the wire labels of the input wires, W_0 , W_1 , W_2 and W_3 . Recall that Alice's input to the 2PC protocol are $x = x_0x_1$ where $x_1, x_0 \in \{0, 1\}$. Alice communicates her input to Bob, without revealing values of x_i , by sending wire labels $W_0^{x_0}$ and $W_1^{x_1}$. Bob does not know the values of x_0 or x_1 , because he is simply receiving two ciphertexts, and does not know which value the ciphertexts represent. The only information that Bob has is the fourth column of the garbled tables, and that does not provide any information about the semantic value of the wire labels.

Recall that Bob's input to the 2PC protocol is $y = y_0y_1$ where $y_1, y_0 \in \{0, 1\}$. Alice now wants to send Bob $W_2^{y_0}$ and $W_3^{y_1}$, but she does not know and cannot know (for the sake of security) y_0 and y_1 . Alice and Bob can achieve this by using Oblivious Transfer, as described in Chapter 1. As an example, we look at wire W_2 . Alice has two possible values that she wants to send to Bob: W_2^0 and W_2^1 . Alice only wants Bob to acquire one of the values, because otherwise he can decrypt multiple rows garbled table. [{AL-16: introduce this OT notation in Chapter 1}](#) Bob wants to receive $W_2^{y_0}$, as that wire label corresponds to his input. So Alice and Bob do $\text{OT}(\{W_2^0, W_2^1\}, y_0)$, so that Alice obviously sends Bob the correct wire label. Alice and Bob also perform OT on wire W_3 ; in particular, they do $\text{OT}(\{W_3^0, W_3^1\}, y_1)$.

Alice is finished garbling the circuit, and communicating information to Bob. Bob

has everything he needs to ungarble, or decrypt, the circuit. Bob starts with gate G_0 , for which he has the fourth column of T_{G_0} , a wire label for wire 0, which I denote W_0^* since Bob does not know which wire label it is, and W_1^* accordingly. Bob starts with the first row of T_{G_0} and tries decrypting the value using W_0^* and W_1^* . Formally, Bob tries $\text{Enc}_{W_0^*, W_1^*}^{-1}(T_{G_0}[0])$ where $T_{G_0}[0]$ represents the value in the zeroth row of T_{G_0} . Bob tries decrypting the values in all four rows of the garbled table T_{G_0} , but only one should work, since the other decryptions use the incorrect keys. For Bob to recognize that encryption is failing, we add an additional property to the encryption scheme: the output of the decryption function should indicate whether the decryption was valid. A single additional bit, where 0 represents that decryption failed and 1 represents that decryption was successful. It is noteworthy that such a property is common to encryption schemes, and can be added to any existing encryption if necessary. With this property, as Bob tries decrypting all four rows of the garbled table, only one should decrypt correctly. Because of the way Alice constructed the garbled table, Bob knows that this correctly decrypted value is one of the wire labels of W_5 , the output wire of gate G_0 . Bob then assigns this decrypted value to W_5^* , and uses the value as the input wire when decrypting gate G_2 .

Bob repeats this same process for gate G_1 and for gate G_2 . For gate G_2 , Bob uses wire labels W_5^* and W_6^* which he acquired by ungarbling gates G_0 and G_1 . Bob notifies Alice after acquiring W_7^* . Alice sends him values W_7^0 and W_7^1 . If $W_7^* = W_7^0$, then the function output is 0 and Bob notifies Alice that the output was 0. Otherwise if $W_7^* = W_7^1$, then the function output is 1 and Bob notifies Alice that the output was 1.

Bob and Alice have now securely computed the function.

[{AL-17: Match notation in this description with the algorithms}](#)

2.3.1 Security of Garbled Circuits

We now discuss the security of garbled circuits without proof.

Algorithm 1 Garble Circuit

Require: Circuit $f(x, \cdot)$ **Ensure:** Populate garbled tables $f(x, \cdot).tables$. **for** wire w_i in $f(x, \cdot).wires$ **do** Generate two encryption keys, called garbled values, W_i^0 and W_i^1 . Assign (W_i^0, W_i^1) to w_i . **end for** **for** gate g in $f(x, \cdot).gates$ **do** Let w_i be g 's first input wire. Let w_j be g 's second input wire. Let w_k be g 's output wire. **for** $(u, v) \in \{(0, 0), (0, 1), (1, 0), (1, 1)\}$ **do** $T_g[u, v] = \text{Enc}_{W_i^u}(\text{Enc}_{W_j^v}(W_k^{g(u,v)}))$ {AL-18: use DKC} **end for** $f(x, \cdot).tables[g] = T_g$. **end for**

Algorithm 2 Evaluate Circuit

Require: $(input_wires, tables, gates)$ **for** Input wire w_i in $input_wires$ **do** \triangleright retrieve garbled values of input wires Perform $\text{OT}(w_i, x_i)$ \triangleright retrieve $W_i^{x_i}$ from Alice Save the value to w_i . **end for** **for** Gate g in $gates$ **do** \triangleright compute the output of each gate. Let w_i be g 's first input wire. Let w_j be g 's second input wire. Let w_k be g 's output wire. **Require** w_i and w_j have been assigned garbled values. **Require** w_k has not been assigned a garbled value. **for** $(u, v) \in \{(0, 0), (0, 1), (1, 0), (1, 1)\}$ **do** {AL-19: use DKC} $temp = \text{Dec}_{w_j}(\text{Dec}_{w_i}(tables[g][u, v]))$ **if** $temp$ decrypted correctly **then** $w_k = temp$ **end if** **end for** **end for**

The definition of security of a 2PC protocol Π given early in Chapter 2 asked us to compare the execution of Π in the real world to an ideal 2PC execution with a trusted third party Carlo. This definition is extremely useful; however, it is difficult to adapt it to intuiting the security of a given protocol. To think about the security of garbled circuits, we think about the information that Alice and Bob.

The only information that Alice receives from Bob during the execution of garbled circuits is in the OT stage. Thereby the security on Alice's side is dependent on the security of OT, whose security we describe in Chapter 1. Hence, we can confidently say that Alice does not learn anything during garbled circuits.

The security of Bob is more complicated. Bob learns three sets of information: a single wire label for each wire (i.e. input labels) and the fourth column of the garbled table for each gate. The input labels are acquired in two ways: some are sent naively by Alice and some are acquired via OT. We can be confident that Bob doesn't learn any extra information in the process of receiving the labels, since we are confident that OT does not reveal information, and other labels are coldly sent by Alice. Moreover, knowing input labels associated with Alice's input doesn't give Bob any information, since he cannot tell if the label is associated with 0 or 1.

The fourth column of the garbled table by itself does not reveal any information to Bob, as it is simply the encryption of things. Bob has the keys to decrypt some of the values in the garbled table. We need to be sure that Bob only has the keys to decrypt the single, intended row of the table. Bob can only decrypt a row of the table if he has both of the labels used as keys. For gates on that are connected to input wires, Bob only has a single label for each wire, therefore he must only be able to decrypt a single row. For subsequent gates, Bob will only ever have a single label for each wire, meaning he can only decrypt a single row.

For Bob to learn extra information, he needs to acquire two wire labels for a single wire. If he can acquire two labels, then he can decrypt the circuit using that wire

label, and with some thinking learn about Alice's input. Based on the information that Bob has, there is no way he can access two wire labels for a single wire. And that is the intuition behind the security of garbled circuits. Of course without a formal proof using computational indistinguishability, we cannot be confident that garbled circuits are secure, as there may be an attack that eludes us.

2.3.2 Notes about complexity

There are three things to think about when considering the complexity of garbled circuits. The first is the amount of information that needs to be communicated per gate. Garbled circuits require 4 ciphertexts, that is 4λ bits. On top of this communication is Alice sending her input labels, and the communication required to complete OT for Bob's input labels. OT, if used naively, is a significant contributor to overall bandwidth.

The second thing to consider is the amount of computation that Alice performs. Alice performs 4 encryptions with a dual-key cipher for each gate.

The third cost to consider is the amount of computation that Bob performs. Bob performs 4 decrypts with the dual-key cipher for each gate.

These constraints are all important, but in practice the biggest bottlenecks are the communication per gate and Bob's computation, which are correlated as we see in Chapter 3. A circuit that computes AES requires approximately 35,000 gates. If 4 ciphertexts per gate need to be communicated, and AES-128 is the encryption scheme, then $4 * 128 * 35,000 \text{ bits} = 2.24 \text{ megabytes}$: a huge amount of communication for just encryption!

2.4 GMW

Where Yao's protocol is premised on encrypting gates individually, GMW's protocol for garbling circuits is premised on secret sharing and performing operations on the shared secrets. In general, secret sharing is a class of methods for distributing a secret to a group of participants, where each participant is allocated a *share* of the secret. The secret can only be reconstructed when a sufficient number of the participants combine their shares, but any pool of insufficient shares yields no information about the secret.

GMW begins by having Alice and Bob secret share their inputs, so each party now has a collection of *shares*. Algorithm ?? describes this process in more detail. Then Alice and Bob perform a series of operations on their shares, which are dictated by the gates in the function they wish to compute. As with Yao's protocol, a gate may either compute XOR, AND or NOT. Each operation requires a different series of operations, which are described in Algorithm ?. Finally, Alice and Bob publicize their shares to each other, at which point each party will have sufficient shares to compute the output of the function.

Chapter 3

Improving MPC

{AL-20: cite Mike Rosulek's talk and website} {AL-21: Mike's website has all the sources...} {AL-22: figure out why GRR3 is secure} {AL-23: Ask about notation: A_0 verse W_i^0 }

The aim of this chapter is to outline new techniques that improve the performance of garbled circuits and discuss their costs and benefits.

We consider three metrics to understand the costs benefits of the new techniques: size of the garbled table, garbler-side computation and evaluator-side computation¹. In the classic scheme of garbled circuits, the garbler and evaluator communicate wire labels of input wires and the fourth column of the garbled table for each gate. Communicating the input labels is an unavoidable cost; this information needs to be exchanged. However, many of the input labels are exchanged with oblivious transfer, discussed in chapter 1, which has been improved. We discuss the improvement to oblivious transfer at the end of this chapter.

New techniques also reduce the size of the garbled table, the other information exchanged between the garbler and evaluator.² In the classical garbled circuit

¹In this chapter, the term garbled table is used interchangeably with the fourth column of the garbled table.

²In chapter 2, Alice was the garbler and Bob was the evaluator. For clarity in chapter 3 and beyond, the term *garbler* to allude to the person who garbles the circuit, and the term *evaluator* to

scheme, the evaluator sends all four rows of the garble table to the evaluator. New techniques, wherein the evaluator samples wire labels *intelligently*, reduce the number of rows of the garbled table that are communicated. Reducing the size of the garbled table is important, because it is a measure of the bandwidth requirement of the scheme. The garbled table amount to most of the communication, and the number of garbled tables scales with the size of the circuit - that is, a larger circuit means more gates means more garbled tables. Moreover, bandwidth is the most important factor in assessing the performance of a garbled circuit scheme. This is because communication of the internet is slower than local computation.

We also consider garbler-side computation, which is the number of encryptions that the garbler performs in preparing a garbled table. In the classical garbled circuit scheme, the garbler performs four encryptions; however, this number is not often reduced. Typically, the new techniques increase or alter the computation that the garbler performs.

Finally, we consider evaluator-side computation, which is the number of decryptions the evaluator performs in ungarbling a gate - that is, finding the output wire label. In the Classical garbled circuit scheme, the evaluator trial decrypts each row until one of the decryptions is successful. Hence, on average, the evaluator performs 2.5 decryptions. New techniques reduce the evaluator-side computation to require a single decryption, a good performance improvement.

Table 3 shows an overview of the cost of all improvements made to garble circuits. The table is split into three sections: size, eval cost and garble cost, corresponding to three metrics mentioned above. Size is the number of rows in the garbled table. Eval cost is number of encryptions that the evaluator performs per gate, and decryptions is the number of decryptions that the garbler performs per gate. Because many of the techniques have different effects on XOR gates and AND gates, each section is divided

describe the person who evaluates, or decrypts, the garbled circuit.

Garbled Circuit Improvement	Size ($x\lambda$)		Eval Cost		Garble Cost		Assumption
	XOR	AND	XOR	AND	XOR	AND	
Classical	4	4	1024	μ s	a	b	a
Point and Permute	TBD	TBD	TBD	μ s	a	b	a
GRR3	TBD	TBD	TBD	μ s	a	b	a
Free XOR	TBD	TBD	TBD	μ	a	bs	a
GRR2 XOR	TBD	TBD	TBD	μ	a	bs	a
FleXOR	TBD	TBD	TBD	μ	a	bs	a
Half Gates	2	0	2	0	4	0	a

Table 3.1: Summary of Garbled Circuit Improvements. GRR3 stands for garbled row reduction 3 and GRR2 stands for garbled row reduction 2. {AL-25: two halves make a whole paper}

into two columns showing the metrics for AND gates and XOR gates separately.

The goal of this chapter is primarily to explain each row in this table. We will understand the technique, and its associated computational and bandwidth costs. We start with the earliest technique and move chronological to the most recent improvement.

3.1 Point and Permute

{AL-26: add to chapter 2 that rows need to be permuted!}

The *point and permute* technique speeds up the evaluator's computation of the garbled table by removing the need to trial decrypt the ciphertexts; instead, the garbler subtly communicates which row of the garbled table to decrypt. Recall that the garbled table is randomly permuted - that is, garbler randomly reorders the rows of the garbled table before sending the garbled table to the evaluator³. Upon receiving the circuit, the evaluator trial decrypts each row of the garbled table until a decryption succeeds⁴.

Point and permute enables the evaluator to bypass the trial decryption step, so

³This is required for security. See chapter 2 for more information

⁴Recall that the decryption algorithm outputs a single bit indicating whether or not the decryption was successful. For more information, see chapter 1.

that they decrypt the correct wire label the first time. When garbling a gate, the garbler randomly assigns a select bit 0 or 1 to each wire label of the gate's input wires, where wire labels of the same wire have opposite bits. More formally, b is randomly sampled from $\{0, 1\}$. Then A_0 gets select bit b and A_1 gets select bit $1 - b$. The garbler permutes the garbled table based on the select bits, and appends the select bits each wire label. When the evaluator ungarbles the gate, they use the select bits appended to each of the two input wire labels to determine which row to decrypt. For example, if the select bits on the end of the input wire labels are 1 and 0, then the evaluator decrypts the third row.

Tables 3.1 show an example of point and permute. A, B and C are wires, where A_0 and A_1 are the zeroth and first wire label of wire A respectively. The left table shows the select bits of the input wires A and B . The garbler gives A_0 select bit 0, determining that A_1 has select bit 1. Likewise, the garbler gives B_0 select bit 1, determining that B_1 has select bit 0.

The garbler then permutes the garbled table based on the select bits. The permuted table is shown on the left in table 3.1. When evaluating this gate, the evaluator has A_* and B_* with select bits a and b . The evaluator decrypts the ciphertext in the row corresponding a and b .

Intuitively, point and permute is secure because the select bits are independent of the truth value (also known as semantic value) of a the wire. Thus it is secure for the garbler to permute the table based on the select bits, and it secure for the garbler to send the select bits to th evaluator.

Point and permute slightly increases garbler-side computation to substantially decrease evaluator-side computation. The garbler samples 4 additional random bits, and the evaluator performs a single decryption. Without point and permute, the evaluator needs to decrypt 2.5 ciphertexts on average, hence the garbler performs roughly 1.5 fewer decryptions per gate. The overall bandwidth is increased by 4 bits

Select Bit	Wire Label	Select Bits	Encryption
0	A_0	(0,0)	$\text{Enc}_{A_0,B_1}(C_1)$
1	A_1	(0,1)	$\text{Enc}_{A_0,B_0}(C_0)$
1	B_0	(1,0)	$\text{Enc}_{A_1,B_1}(C_0)$
0	B_1	(1,1)	$\text{Enc}_{A_1,B_0}(C_0)$

Table 3.2: Garbled Gate for Point and Permute

per gate: a tiny constant increase.⁵

{AL-27: Is this not an XOR/And gate? Make it one of them}

3.2 Garbled Row Reduction 3

{AL-28: check notation for get/sample in figure. \leftarrow is used for both} Garbled Row Reduction 3 (GRR3) reduces the size of the garbled table from 4 ciphertexts to three 3 ciphertexts. In a classical garbled circuit scheme, the wire labels for each wire are chosen prior to garbling any gates. In GRR3, the garbler samples values for the input wire labels prior to garbling, and the garbler gives all other wire labels values as they generate each garbled table.

Suppose the garbler is garbling an XOR gate with input wires A and B and output wire C . The garbler begins by using the point and permute method, sampling select bits and permuting the garbled table. In GRR3, the garbler then set the ciphertext in the top row of the garbled table equal to a value that decrypts to 0^n , the string of n zeros. Specifically, C_* , the wire label on the top row, is set to $\text{Enc}_{A_*,B_*}^{-1}(0^n)$. The garbler sends the bottom three rows of the garbled table to the evaluator. When evaluating the garbled gate, if the evaluator sees that the select bits of the input wires indicate to decrypt the first row, then the evaluator simply assumes the row to have value 0^n and decrypts it as usual. In that case, the garbler set C_* to $\text{Enc}_{A_*,B_*}^{-1}(0^n)$. Otherwise if the select bits indicate to decrypt the second, third or fourth row, the garbler decrypts the row as per usual.

⁵The value is constant in the sense that it is independent of the security parameter.

Select Bit	Wire Label	Select Bits	Encryption
0	A_0	(0,1)	$\text{Enc}_{A_0,B_0}(C_0)$
1	A_1	(1,0)	$\text{Enc}_{A_1,B_1}(C_0)$
1	B_0	(1,1)	$\text{Enc}_{A_1,B_0}(C_0)$
0	B_1		

$C_0 \leftarrow \{0,1\}^n$
$C_1 \leftarrow \text{Enc}_{A_0,B_1}^{-1}(0^n)$

Table 3.3: Example garbled gate using point and permute and garbled row reduction 3. The gate being computed is given in figure **Make it 23:30 in mike's talk**

Tables 3.3 gives an example of garbling an XOR gate. The left table shows the select bits of wire labels A_0, A_1, B_0 and B_1 . The right table shows the garbled table, in which the top row, the row associated with select bits (0,0), is missing, as the row is assumed to have value 0^n . The bottom table shows the values of C_0 and C_1 . The value of C_0 is randomly sampled from $\{0,1\}^n$.

In considering the security of GRR3, we consider the effect of always setting one of the wire labels, C_* , to 0^n . Since the evaluator does not know whether C_0 or C_1 is set to 0^n , the evaluator does not learn any information about the semantic representation of C_* .

GRR3 offers good performance benefits. Garbler-side computation is the same, except that a decryption is performed in place of an encryption in constructing the first row of the garbled table. Evaluator-side computation is the same as point and permute: the evaluator performs a single decryption. Finally, GRR3 reduces the size of the garbled table from 4 ciphertexts to 3 ciphertexts, contributing a 25% reduction in bandwidth.

3.3 Free XOR

The Free XOR technique makes the computation of XOR gates free, in the sense that no garbled table need to be communicated. The evaluator can compute C_* from only A_* and B_* .

Like GRR3, the free xor techniques takes advantage of carefully crafted wire labels, even input wires. To start, the garbler randomly samples a ciphertext Δ from $\{0, 1\}^n$. For each input wire A , let A_0 be randomly sampled from $\{0, 1\}^n$ as before, and set $A_1 = A_0 \oplus \Delta$.

If the garbler is garbling an XOR gate, then the garbler does not constructed a garbled table or use select bits. The garlber sets $C_0 = A_0 \oplus B_0$ and sets $C_1 = C_0 \oplus \Delta$.

The evaluator, when evaluating an XOR gate, simply computes $C_* = A_* \oplus B_*$. As simple as it is, the evaluator will always acquire the correct value for C_* based on the semantic value of A_* and B_* . The math for each of the four cases is shown:

$$A_0 \oplus B_0 = C_0 \tag{3.1}$$

$$A_0 \oplus B_1 = A_0 \oplus (B_0 \oplus \Delta) = (A_0 \oplus B_0) \oplus \Delta = C_1 \tag{3.2}$$

$$A_1 \oplus B_0 = (A_0 \oplus \Delta) \oplus B_0 = (A_0 \oplus B_0) \oplus \Delta = C_1 \tag{3.3}$$

$$A_1 \oplus B_1 = (A_0 \oplus \Delta) \oplus (B_0 \oplus \Delta) = (A_0 \oplus B_0) = C_0 \tag{3.4}$$

For any wire A , since A_1 is dependent on A_0 , we often simplify notation such that the wire labels for A are A and $A \oplus \Delta$, omitting the subscript.

Free XOR is compatible with point and permute and GRR3; however, since XOR does not require a garbled table, GRR3 is only used on AND gates.

One interesting implication of using the Free XOR technique is that an added assumption must be made to our encryption algorithm. Since Δ is part of the key and part of the the payload⁶ of the encryption algorithm, the encryption algorithm must be secure under the circularity assumption. Fortunately, the popular encryption scheme AES-128 is presumed to be secure under the circularity assumption.

Free XOR dramatically reduces bandwidth. But since XOR gates are much cheaper than AND gates, circuits with more XOR gates perform faster. Many pro-

⁶The payload is the value that is being encrypted

grams have been made to optimize the number of XOR gates and minimize the number of AND gates (while minimizing the size of the entire circuit of course). The garbler-side computation is reduced: constructing the xor garbled table does not require 3 encryptions and 1 decryption. Evaluator-side computation is reduced likewise: xor gates do not require any decryption.

3.4 Garbled Row Reduction 2

{AL-29: <https://web.engr.oregonstate.edu/~rosulekm/scbib/index.php?n=Paper.PSSW09>}

Garbled row reduction 2 (GRR2) reduces the size of the garbled table of AND gates to 2 ciphertexts. Unfortunatley, GRR2 is not compatible with Free XOR, and GRR2 is less efficient than Free XOR combined with GRR3 for most circuits, so GRR2 is not often used by itself in practice.

Suppose the garbler is garbling an AND gate with input wires A and B and output wire C . For all $a, b \in \{0, 1\}$, let $V_{a,b} = H(A_a, B_b)$. Since the evaluator acquires A_a and B_b for specific a and b , they are also able to acquire $V_{a,b}$ for the same specific a and b . The garbler constructs a polynomial P to be the unique 2-degree polynomial passing through the points $(1, V_{0,0})$, $(2, V_{0,1})$ and $(3, V_{0,1})$ - these V 's are selected because the garbler is garbling an AND gate. And they set C_0 to $P(0)$. The garbler next constructs a second polynomial Q to be the unique 2-degree polynomial passing through points $(4, V_{1,1})$, $(5, P(5))$ and $(6, P(6))$. The garbler then sets C_1 to $Q(0)$. The garbler sends $(5, P(5))$ and $(6, P(6))$ to the evaluator.

When de-garbling, the evaluator holds a, b, A_a, B_b for some $a, b \in \{0, 1\}$, $(5, P(5))$ and $(6, P(6))$. With this information, the evaluator constructs polynomial R to be unique the degree-2 polynomial passing through points $(2a+b+1, H(A_a, B_b))$, $(5, P(5))$ and $(6, P(6))$. Polynomial R is either P or Q , depending on the values of a and b . The evaluator simply sets C_* to be $R(0)$.

GRR2 is incompatible with Free XOR since it sets C_0 and C_1 implicitly and unpredictably. Free XOR requires that $C_0 = C_1 \oplus \Delta$ for some global delta, which is not achievable with the randomly constructed polynomials.

{AL-30: discuss security. Why is free xor secure?} GRR2 alters garbler-side computation by having the garbler construct two polynomials instead of encrypting ciphertexts. This is slightly more computation than required by FreeXOR and GRR3, but not an undermining amount. Evaluator side computation increases slightly, as the evaluator constructs a polynomial instead of decrypting two ciphertexts.

3.5 FleXOR

{AL-31: passive tone} After the creation of GRR2, secure computation was at an awkward point. Circuits with many XOR gates were computed most quickly with Free XOR and GRR3, but circuits with many AND gates were computed most quickly with GRR2. FleXOR reconciles GRR2 with Free XOR, resulting in a scheme that is universally faster.

Recall that GRR2 is incompatible with Free XOR because the wire labels are uncontrollably created by random polynomials, where Free XOR requires that $C_1 = C_0 \oplus \Delta$ for some global delta. FleXOR solves this problem in a straightforward fashion: correct the delta value of output wires of AND gates such that the output wires use the global delta. To correct the value, FleXOR adds a unary gate after each AND gate that corrects the wire label. ⁷

Suppose an XOR gate has input wires A and B and output wire C . Input wires A and B each come from an AND gate, so their labels are the result of the polynomial interpolation of GRR2. A has labels A and $A \oplus \Delta_A$ and B has labels B and $B \oplus \Delta_B$. To perform Free XOR, A , B and C need to be using the same delta value. The garbler

⁷A unary gate is a gate that takes a single input wire and outputs a single wire. The unary gate does not change the semantic meaning of a wire label - that is, whether it represents 0 or 1. The unary gate merely alters the actual value of the wire label or ciphertext.

adds an extra gate between A and B and the XOR gate that adjusts their XOR value to the correct value. A and $A \oplus \Delta_A$ change to A' and $A' \oplus \Delta$. B and $B \oplus \Delta_B$ change to B' and $B' \oplus \Delta$. Since A, B and C have the same delta value, FreeXOR is used.

The unary gate maps $A, A \oplus \Delta_1 \rightarrow A', A' \oplus \Delta$, where Δ is the correct delta value for the XOR gate.

FleXOR is made more efficient by not correcting the outputwire of every AND gate. For example, if an output wire of an AND gate is immediately inputted into another AND gate, the wire label does not need to be corrected. Moreover, the wire labels do not even need to be corrected to a global delta. Free XOR only requires that the three wires involved in the XOR gate, A, B and C , use the same delta, so each XOR gate has its own Δ . For example, A and C may share a delta but B may have a different delta, so it is sufficient to only correct B 's delta value.

FleXOR is fastest when AND gates are grouped together and XOR gates are grouped together, since fewer unary gates will be required. Thereby, FleXOR creates an optimization problem: place XOR gates, AND gates and unary gates to minimize bandwidth requirements. However, FleXOR's optimization problem is too computationally expensive, but on the up side, analysis reveals that FleXOR requires on average an extra 0 or 1 ciphertext per gate, a relatively small amount. That cost comes at the benefit of a Free XOR circuit, and 1 fewer ciphertext for each AND gate.

FleXOR requires slightly more garbler-side computation than FreeXOR and GRR2, since the garbler must create the unary gates. The evaluator-side computation is the same as FreeXOR and GRR2, with the additional computation of the unary gates, which is small. The size of the XOR garbled table is 0, and the size of the AND garbled table is 2. But there is the additional communication of the unary gates' garbled tables. The garbled table of the unary gate has 2 ciphertexts, and the number of unary gates depends on the circuit.

FleXOR is intuitively secure, since the only additional information beyond GRR2 and FreeXOR is the unary gates. The unary gate is secure, since it functions the same as a normal garbled gate except that an input is missing. However, FreeXOR requires the circularity assumption, The circularity assumption is discussed in the section on FreeXOR.

3.6 Half Gates

{AL-32: for some reason it uses hashes instead of encryption in the paper. not sure why. and in the talk, mike uses enc. in our code, alex m used hash} Half Gates is the most recent improvement to garbled circuits. The goal of Half Gates is to make AND gates cost two ciphertexts, while preserving properties necessary for Free XOR without adding unary gates. Half gates work by splitting an AND gate into two half gates. Each half gate is an AND gate, where one of the inputs is *known* to a party. One half gate is the garbler-half-gate, where the garbler *knows* the semantic value of one of the input wires. Similarly, the evaluator knows the semantic value of one of the input wires of the evaluator-half-gate. We then combine the garbler-half-gate and the evaluator-half-gate to produce a generic AND gate.

{AL-33: where do the hashes come from?} We first examine the garbler-half-gate. Imagine an AND gate with input wires A and B and output wire C , where the garbler knows the semantic value of A . This means that the garbler knows whether A_* is A_0 or A_1 . Let a, b and $c \in \{0, 1\}$ represent the semantic value of A, B and C respectively. If $a = 0$, then $c = a \wedge b = 0$, so the garbler creates a unary gate that always outputs C_0 . If $a = 1$, then $c = a \wedge b = b$, so the garbler creates a unary identity gate that outputs C_b . We combine the unary gates into a single garbled table, shown in table 3.6. Since the evaluator has B_* , the evaluator can compute the label of the output wire of the AND gate by xoring the rows in the garbled table by $H(B)$ or $H(B \oplus \Delta)$

Garbled Table for $a = 0$	Garbled Table for $a = 1$	→	Garbled Table for any a
$\text{Enc}_{B_0}(C_0)$	$\text{Enc}_{B_0}(C_0)$ $\text{Enc}_{B_1}(C_1)$		$\text{Enc}_{B_0}(C_0)$ $\text{Enc}_{B_1}(C_0 \oplus a\Delta)$

Table 3.4: Left: garbler-half-gate garbled table for $a = 0$. Middle: garbler-half-gate garbled table for $a = 1$. Right: garbler-half-gate garbled table for arbitrary a .

to get C or $C \oplus \Delta$ respectively.

Using point and permute and the GRR3 trick reduces the size of the garbler-half-gate garbled table to 1 ciphertext. The garbler chooses C_0 such that the first of the top row of the garbled table is the all zeros ciphertext, and therefore is not sent to the evaluator. If the evaluator should decrypt the ciphertext on the top row (as directed by point and permute), then the evaluator assumes the ciphertext to be all 0s.

The evaluator-half-gate is somewhat different from the garbler-half-gate. Consider an AND gate where the evaluator *knows* the semantic value of a . If $a = 0$, then the evaluator should acquire C_0 . Otherwise if $a = 1$, then the evaluator should acquire $C_0 \oplus b\Delta$. In this case, it suffices for the evaluator to obtain $\Omega = C_0 \oplus B_0$, as the evaluator simply sets C_* to $B_* \oplus \Omega$ which sets C_* equal to $C_0 \oplus b\Delta$.

Table 3.6 the garbled table for the evaluator-half-gate. This table does not need to be permuted, since the evaluator knows a . Again, we use the GRR3 trick, and set C_0 to $\text{Enc}_{A_0}^{-1}(0^n)$, eliminating the need to send the top row of the garbled table.

Garbled Table for any a
$\text{Enc}_{A_0}(C)$ $\text{Enc}_{A_1}(C \oplus B)$

Table 3.5: Evaluator-half-gate garbled table.

We now put the two half gates together to form an AND gate. Consider the following where r is a random bit generated by the garbler:

$$a \wedge b = a \wedge (r \oplus r \oplus b) \quad (3.5)$$

$$= (a \wedge r) \oplus (a \wedge (r \oplus b)). \quad (3.6)$$

The first AND gate, $a \wedge r$, can be computed with a garbler-half-gate - the garbler *knows* r . Furthermore, if we can let the evaluator know the value of $r \oplus b$, then the second AND gate, $(a \wedge (r \oplus b))$, can be computed with an evaluator-half-gate - the evaluator *knows* $r \oplus b$. And the XOR gate can be computed with free xor at the cost of no ciphertexts.

It is secure for the garbler to give $r \oplus b$ to the evaluator, since r is random so it blinds the value of b . The value of $r \oplus b$ can be communicated to the evaluator for free: use the select bit (from the point and permute technique) of the false wire label on wire b (so r is the select bit on the true wire label of wire b).

The overall cost of using Half Gates for AND gates is four encryptions for the garbler, two decryptions for the evaluator, and the communication of 2 ciphertexts. Half Gates guarantees only two ciphertexts are needed per AND gate, but the trade-off is the additional computation for both parties. With FleXOR, the number of ciphertexts that need to be communicated may vary, but there is less computation required.

3.7 Improving Oblivious Transfer

[{AL-34: Fix this up}](#) This section discusses improvements to oblivious transfer, the method by which the garbler communicates the wire labels corresponding to the evaluator's inputs to the evaluator. At its most basic, oblivious transfer enables Alice to potentially send one of two messages to Bob. Bob selects which message he wants to receive. The desirable security properties are (1) Alice does not know which message she sent to Bob and (2) Bob cannot infer any information about the message that he did not receive.

Oblivious transfer is a key component in secure computation, and often consumes a substantial portion of the secure computation protocol's time. There are two major

improvements to oblivious transfer. The first is called *OT-extension*. When using OT with a garbled circuit scheme, Alice and Bob are not exchanging just one message: they exchange a wire label for each of Bob's inputs. This means that Alice and Bob are performing a linear number of OTs based on the number of inputs. OT-extension reduces the entire OT phase to a constant number of OT operations. In particular, Alice and Bob run a constant number of OT in order to generate a polynomial number of exchanged messages.

The second improvement to OT is called *OT-preprocessing*. In OT-preprocessing, Alice and Bob performs the expensive OT operation ahead of time, in what is called the offline phase, on random values. Then, during the online phase, Alice and Bob quickly exchange values to correct the pre-shared messages. Alice and Bob only exchange a single message during the online phase, and there is no computationally expensive math required like there is during the real OT operation.

In our work with secure computation, OT-preprocessing substantially improved the performance of garbled circuits, and OT-extension was less useful. If a garbled circuit scheme uses an offline phase and thereby uses OT-preprocessing, then OT-extension is largely unnecessary. The parties take the extra needed to send all of the wire labels in the offline phase, since the time does not matter. What we the parties are truly concerned about is online time. During online time, OT-preprocessing offers huge advantages.

Chapter 4

Chaining Garbled Circuits

In this chapter we introduce component-based garbled circuits. Component-based garbled circuits allow for most of the work involved in building and communicating a garbled circuit to occur in an offline phase before the inputs or function to compute are known. This offers a significant improvement to the online performance of garbled circuits.

We begin with the observation that most functions are composed of many smaller components. For example, many statistical functions are a composition of matrix and arithmetic operations. The gist of component-based garbled circuits is then to build and communicate many garbled components in the offline stage. Then, in the online stage, the garbler and evaluator choose their inputs and a function to compute. The garbler communicates input labels and auxiliary information, to be discussed, to the evaluator, who stitches together the garbled components into the larger function.

The idea may make more sense with a real world example. Imagine that two banks integrate secure computation into their daily transactions. At night when activity is low, the banks' servers exchange many garbled circuits, and then during the day, they use the pre-exchanged garbled circuits to quickly perform secure computation. The computational requirements for the banks, when they compute the function, is

only to exchange input labels and for the evaluator-bank to evaluate the garbled circuit, while, importantly, the banks preserve the ability to choose their inputs and the function to be computed at the time of the computation.

4.1 Chaining

In this section we describe how to chain together garbled circuits.

Suppose that Alice and Bob are chaining two garbled circuits. During an offline phase, Alice garbles many circuits and sends them Bob. Later, during the online phase, Alice sends Bob input labels, some via oblivious transfer, and then Alice sends Bob *link* labels. A link label converts an output wire from one garbled circuit to an input wire of another garbled circuit. Say that Alice and Bob want to set output wire W_i of a garbled circuit to the input wire W_j of another garbled circuit. We suppose that Alice and Bob are using Free XOR, where each garbled circuit uses the same global Δ .

We formally describe component-based garbled circuits as a tuple of three algorithms (**Garble**, **Link**, **Eval**). The **Garble** algorithm is the same algorithm as before; it takes as inputs a circuit C and outputs a garbled circuit GC . We extend the **Garble** algorithm to also take a component-circuit c as input, instead of a complete circuit C , which outputs a garbled component GC_c , input wire set e_c and output wire set d_c .

The **Link** algorithm is new and designed for component-based garbled circuits. **Link** takes as input two garbled components $c_0 = (GC_0, e_0, d_0)$ and $c_1 = (GC_1, e_1, d_1)$, and a mapping of output wires of c_0 to input wires of c_1 . **Link** outputs the *link* labels needed to convert the output wires c_0 to input wires e_1 . Suppose that output wire $W_i \in d_0$ has labels (W_i^0, W_i^1) , and input wire $W_j \in e_1$ has wire labels (W_j^0, W_j^1) . To link W_i to W_j , **Link** outputs $W_{ij} = W_i^0 \oplus W_j^1$. Since Free-XOR is being used, we have that $W_i^0 \oplus W_i^1 = W_j^0 \oplus W_j^1 = \Delta$ for some random value Δ . W_{ij} is sufficient

to transform a wire label of W_i to a wire label of W_j while preserving their semantic value. {AL-35: Do we need more explanation of how chaining works here? I'm guessing another paragraph explaining how the Delta moves over}

The **Eval** algorithm is based on the original **Eval** algorithm used in normal garbled circuits, but with chaining added where needed. The **Eval** algorithm takes as input a list of a garbled components and $\{c_i\}$, linking labels $\{W_{ij}\}$, and it outputs output labels $\{Y_i\}$. {AL-36: Why does it not take input labels as input also}. **Eval** starts with the inputs, and then proceeds component by component, evaluating each component in order to get the component output wire label. Where necessary, **Eval** uses component output wire labels and link labels to compute the appropriate input label for later components. Once all components are evaluated, **Eval** recovers the garbled outputs $\{Y_i\}$ from the output components, and uses d for that component to recover the real output y .

Let us think about how the two banks who use secure computation in their daily operations employ the three algorithms. At night, during the offline stage, the garbler-bank generates many circuit components, garbles them with **Garble**, and sends the garbled circuits to the other evaluator-bank. The garbler-bank and evaluator-bank also perform the offline phase of OT-preprocessing. During the day, when the banks decide that they need to compute some function f , they check that they have sufficient components to compute f . If they do, the garbler-bank runs **Link** on the components used in f , and sends the output link-labels to the evaluator. The garbler and evaluator also exchange wire labels, some via online OT-preprocessing. Now that the garbler has the garbled components, link labels and input labels, they run **Eval**, recovering the output of the function.

4.2 Security of Chaining

In this section, we show how to adapt the standard definition of privacy, presented in chapter 2, to component-based garbled circuits.

In this section we explain why chaining is secure. Informally, chaining is secure since the chaining values do not represent any semantic information, and the evaluator cannot infer any further information from them. The evaluator cannot XOR a chaining value with any other value and gain meaningful information. Say the evaluator has chaining value W_{ij} , chaining W_i to W_j . The evaluator learns nothing by xoring W_{ij} with W_i or W_j , and W_{ij} is not associated with W_k for all $k \neq i, j$, so the evaluator stands to learn nothing.

To think about the security of chaining formally, we modify the experiment of classical garbled circuits to include the chaining step¹. We modify the experiment to the following:

4.3 Single Communication Multiple Connections

Single Communication Multiple Connections (SCMC) is a method of reducing the number of linking values that Alice and Bob exchange. In the naive chaining scheme, online bandwidth scales linearly with the number of wires being chained, since a ciphertext is communicated per wire being chained. For some cases, this is a lot of communication. If Alice and Bob are chaining two garbled circuits that perform matrix multiplication on 100 by 100 matrices with entries that are 16 bits, then Alice sends 8,000 ciphertexts to Bob - 8,000 ciphertexts is 4 kBs. SCMC reduces the chaining operation to require a single ciphertext of communication: an 8,000x improvement.

SCMC works like most garbled circuits improvements: cleverly choosing wire

¹The experiment for classical garbled circuits is given in Chapter 2

labels. The key observation is that consecutive output wires are often chained to consecutive input wires. Suppose we are chaining garbled circuit C_0 with output wires $\{W_i\}_{i \in \text{Outputs}(C_0)}$ to garbled circuit C_1 with input wires $\{Y_i\}_{i \in \text{Inputs}(C_1)}$. Moreover, for all wires W_i being chained in C_0 , W_i is being chained to Y_j where $i = j$; intuitively, a wire in C_0 is being chained to its corresponding wire in C_1 . {AL-37: What does this mean? What is a corresponding wire?} Alice, when garbling, sets all output wires W_i to $W_0 \oplus H(T \oplus (i||b))$, and all output wires Y_i to $Y_0 \oplus H(T \oplus (i||b))$ where H is a hash function.

4.4 Security of SCMC

SCMC is known to be secure in the random oracle model. In the random oracle model, we assume that H is a random function. That is, on any input $x \in \{0, 1\}^*$, H outputs a uniform random bit-string of length λ .

...

Chapter 5

Implementation

I created **CompGC**, a system for chaining garbled circuits. **CompGC** is designed to run a two party garbled circuit protocol from start to finish.

CompGC has a garbler-side and an evaluator-side, which is specified with command line arguments. The **CompGC** system operates in two phases. In the first phase, the offline phase, the garbler sends garbled circuits to the evaluator, and they perform the offline phase of OT-preprocessing. In the second phase, the online phase, the garbler sends input labels for the garbled circuits, some of which are exchanged via the online phase of OT-preprocessing, instructions for how the evaluator should chain the garbled circuits, and output garbled circuits which are garbled circuits composed of a single, unary gate that encode the real output of the function being processed. The following section goes into more detail into the particular of the **CompGC** protocol.

5.1 **CompGC**

I initially constructed **CompGC** on top of **JustGarble**. **JustGarble** is a system for garbling and evaluating circuits, and it only garble circuits and evaluates. **JustGarble** takes as input a circuit and can output a garbled circuit, and similarly, it can take as input a garbled circuit and input wires and return the output of the garbled circuit.

While building **CompGC**, **LibGarble** was being developed by a researcher at University of Maryland. **LibGarble** is a fork of **JustGarble**- that is, the researcher took the code of **JustGarble** as his starting point - and offers the same functionality as **JustGarble**, but it has a better API, is faster, and supports more features, namely half-gates. I eventually started using **LibGarble** in **CompGC**.

Starting with **JustGarble** and **LibGarble**, which perform a very specific task, meant that I needed to add many steps and features to make a fully functional garbled circuit system. I started by adding oblivious Transfer, for which I found an implementation of the Naor-Pinkas algorithm. Naor-Pinkas is basic, semi-honest oblivious transfer. A colleague of mine later added to the oblivious transfer code to support OT-preprocessing.

I will now walk through the 4 basic modes of **CompGC**: garbler-offline, evaluator-offline, garbler-online and evaluator-online. Algorithm 3 outlines the garbler's offline implementation.

Algorithm 3 Garbler Offline

Input: List of circuits, and number of ciphertexts to be OT-preprocessed.

1. Build circuits
 2. Garble circuits (using **LibGarble**)
 3. Assign each garbled circuit an ID
 4. Send garbled circuits and their IDs to evaluator
 5. Save input labels and output labels of garbled circuit to disk
 6. Perform offline phase of OT-preprocessing
 7. Save data from OT-preprocessing to disk
-

Algorithm 4 Evaluator Offline

1. Receive garbled circuits from garbler
 2. Save garbled circuits to disk
 3. Perform offline phase of OT-preprocessing
 4. Save OT-preprocessing data to disk
-

Algorithm 5 Garbler Online

1. Load input labels and output labels of garbled circuits from disk
 2. Load OT-preprocessing data from disk
 3. Load function specification from disk
 4. Generate instructions from function specification
 5. Compute chaining values, and add values to instructions
 7. Perform online stage of OT-preprocessing
 8. Send input wires correspond to garbler's input
 9. Send instructions
-

Algorithm 6 Evaluator Online

1. Load garbled circuits from disk
 2. Load OT-preprocessing data from disk
 3. Perform online stage of oblivious transfer - acquire evaluator's input labels
 4. Receive garbler's input labels
 5. Receive instructions
 6. Following instructions, chaining and evaluating as instructed
-

5.2 LibGarble and JustGarble

{AL-38: fix} CompGC is written on top of LibGarble. LibGarble is a C library that is a fork of JustGarble. JustGarble was produced by BHR in 2012. All it does is create circuits, garble them, and evaluate them in very fast speeds. LibGarble is a rework of JustGarble in which the memory layout and other things are changed to speed up computation.

5.3 Our Implementation: CompGC

CompGC implements the theoretical ideas of naive chaining and SCMC.

5.4 Adding SCMC

5.5 Experimental Setup

talk about what computer was used, bandwidth, network simulation

5.6 Experiments

AES, CBC, Levenshtein

5.7 Results

Chapter 6

Future and Related Work

talk about future and related work here

Conclusion

Here's a conclusion, demonstrating the use of all that manual incrementing and table of contents adding that has to happen if you use the starred form of the chapter command. The deal is, the chapter command in \LaTeX does a lot of things: it increments the chapter counter, it resets the section counter to zero, it puts the name of the chapter into the table of contents and the running headers, and probably some other stuff.

So, if you remove all that stuff because you don't like it to say "Chapter 4: Conclusion", then you have to manually add all the things \LaTeX would normally do for you. Maybe someday we'll write a new chapter macro that doesn't add "Chapter X" to the beginning of every chapter title.

Appendix A

The First Appendix

Appendix B

The Second Appendix, for Fun

References

- Bellare, M., Hoang, V. T., Keelveedhi, S., & Rogaway, P. (2013). Efficient garbling from a fixed-key blockcipher. In *IEEE Symposium of Security and Privacy*, (pp. 478–492).
- Bellare, M., Hoang, V. T., & Rogaway, P. (2012). Foundations of garbled circuits. In *Proceedings of the 2012 ACM Conference on Computer and Communications Security*, (pp. 784–796). ACM.
- Boneh, D. (1998). *Algorithmic Number Theory: Third International Symposium, ANTS-III Portland, Oregon, USA, June 21–25, 1998 Proceedings*, chap. The Decision Diffie-Hellman problem, (pp. 48–63). Berlin, Heidelberg: Springer Berlin Heidelberg. <http://dx.doi.org/10.1007/BFb0054851>
- Goldreich, O. (1995). Foundations of cryptography.
- Goldreich, O., Micali, S., & Wigderson, A. (1987). How to play any mental game. In *Proceedings of the Nineteenth Annual ACM Symposium on Theory of Computing, STOC '87*, (pp. 218–229). New York, NY, USA: ACM. <http://doi.acm.org/10.1145/28395.28420>
- Katz, J., & Lindell, Y. (2007). *Introduction to Modern Cryptography: Principles and Protocols (Chapman & Hall/CRC Cryptography and Network Security Series)*. Chapman and Hall/CRC. <http://www.amazon.com/Introduction-Modern->

[Cryptography-Principles-Protocols/dp/1584885513%3FSubscriptionId%3D0JYN1NVW651KCA56C102%26tag%3Dtechkie-20%26linkCode%3Dxm2%26camp%3D2025%26creative%3D165953%26creativeASIN%3D1584885513](https://www.amazon.com/dp/1584885513?SubscriptionId%3D0JYN1NVW651KCA56C102%26tag%3Dtechkie-20%26linkCode%3Dxm2%26camp%3D2025%26creative%3D165953%26creativeASIN%3D1584885513)

Lindell, Y., & Pinkas, B. (2009). Secure multiparty computation for privacy-preserving data mining. *Journal of Privacy and Confidentiality*, 1(1), 5.

Snyder, P. (????). Yaos garbled circuits: Recent directions and implementations.